

# CRISPR Jamboree 2024: Inference

Gene Katsevich

January 26, 2024

## 1 Overview of datasets, inputs, and outputs

I propose the following for the inference module:

- **Datasets:** We should test on at least one high-MOI and one low-MOI dataset. I propose we try subsets of the [Gasperiini](#) (high-MOI) and [Papalexi](#) (low-MOI) datasets, which are built into the `sceptre` package. These datasets are available [on GitHub](#).
- **Input:** The input will be a `MuData` object with a certain set of minimum required fields (see more below), the sidedness of the test, and any additional arguments to the analysis method.
- **Output:** The output will be an updated `MuData` object, with the  $p$ -values and log fold-change estimates added in the `uns` field.

## 2 MuData structure

We will need to specify precisely how the `MuData` object should be formatted. The current [MuData formatting guidelines](#) do not specify naming conventions precisely enough, and which fields are mandatory versus which fields are optional. Lucas's [sample Gasperiini MuData object](#) is a good starting point, but it has the following limitations:

- It does not include negative control gRNAs.
- It does not include the gRNA-to-cell assignments.
- It does not include the data frame for the results of an analysis method.
- Certain naming choices conflict with the `MuData` formatting guidelines document.

In what follows, I propose an example of a `MuData` object structure that addresses some of the above limitations. To illustrate, I will use a subset of the Gasperiini dataset. Note that this subset is different from Lucas's. For example, it contains non-targeting gRNAs. The `MuData` structure itself is inspired by Lucas's but includes some additions, such as the gRNA-to-cell assignments and the data frame for the results of an analysis method.

```
import mudata as md
import pandas as pd
gasperiini_data = md.read_h5mu("data/gasperini_subset.h5mu")
gasperiini_data

## MuData object with n_obs × n_vars = 44308 × 716
##   obs:   'batch'
##   uns:   'inference_results', 'moi'
##   3 modalities
##     gene:   44308 x 526
##       var:   'symbol'
##     guide:  44308 x 95
##       var:   'targeting', 'guide_target', 'guide_chr', 'guide_start', 'guide_end'
##     guide_assignment: 44308 x 95
```

We can see that this object has three modalities, named `gene` (for gene expression), `guide` (for gRNA expression), and `guide_assignment` (for binary gRNA-to-cell assignments). We need to make sure these names are standardized.

## 2.1 Gene expression modality

The gene expression modality (`gene`) should contain the raw RNA UMI counts; we will leave any normalization to each of the individual analysis methods. The variable names are the ENSEMBL gene IDs:

```
gasperini_data['gene'].var_names[:5].tolist()
```

```
## ['ENSG00000069275', 'ENSG00000117222', 'ENSG00000117266', 'ENSG00000117280', 'ENSG00000133059']
```

The gene names can be optionally be provided in a variable called `symbol`:

```
gasperini_data['gene'].var
```

```
##              symbol
## ENSG00000069275  NUCKS1
## ENSG00000117222  RBBP5
## ENSG00000117266  CDK18
## ENSG00000117280  RAB29
## ENSG00000133059  DSTYK
## ...              ...
## ENSG00000155380  SLC16A1
## ENSG00000196683  TOMM7
## ENSG00000176890  TYMS
## ENSG00000198786  MT-ND5
## ENSG00000198840  MT-ND3
##
## [526 rows x 1 columns]
```

We need to choose a standardized name for the variable storing the gene names; I chose `symbol` because this is what DACC appears to use.

## 2.2 gRNA expression modality

The gRNA expression modality (`guide`) should contain the raw gRNA UMI counts; we will leave any normalization to each of the individual analysis methods. Within the `guide` modality, the variable names are the gRNA IDs:

```
gasperini_data['guide'].var_names[:5].tolist()
```

```
## ['grna_CCGGGCG', 'grna_TGGCGGC', 'grna_AAGGCCG', 'grna_GACGCCG', 'grna_CACACCC']
```

The variables in this modality are similar to the [per-guide metadata format](#) we developed:

```
gasperini_data['guide'].var
```

```
##      targeting  guide_target guide_chr  guide_start  guide_end
## grna_CCGGGCG      1  ENSG00000069482    chr11  68451943.0  68451958.0
## grna_TGGCGGC      1  ENSG00000069482    chr11  68451958.0  68451974.0
## grna_AAGGCCG      1  ENSG00000100316    chr22  39715775.0  39715790.0
## grna_GACGCCG      1  ENSG00000100316    chr22  39715790.0  39715806.0
## grna_CACACCC      1  ENSG00000104131    chr15  44829255.0  44829270.0
## ...              ...              ...              ...
## grna_ATTAGCA      0  non-targeting              -9.0        -9.0
## grna_AGATACC      0  non-targeting              -9.0        -9.0
## grna_ATATGTA      0  non-targeting              -9.0        -9.0
```

```
## grna_GTAGCCT      0    non-targeting      -9.0      -9.0
## grna_TTAGCTT      0    non-targeting      -9.0      -9.0
##
## [95 rows x 5 columns]
```

First, we must specify for each gRNA whether it is **targeting**. (Here this boolean variable shows up as 0/1; probably we should change this.) Next, we must specify for each gRNA what exactly it targets (which gene, which putative enhancer, etc). This is probably similar to **intended\_target\_name** in the per-guide metadata format but I wasn't sure, so I named it **guide\_target** (we should come to a consensus on this). The reason this is required is because we it is most meaningful to test for associations between *targeted elements* and gene expression rather than between *individual gRNAs* and gene expression. Therefore, we must know which guides target the same element. We should have a reserved string for the **guide\_target** for non-targeting guides; I propose **non-targeting**. Optionally, we can specify the genomic coordinates of the region targeted by the guide. We should have standard placeholder values genomic coordinates of non-targeting guides. I propose the empty string for chromosome and -9 for start and end coordinates.

## 2.3 gRNA assignment modality

The gRNA assignment modality (**guide\_assignment**) should contain the binary gRNA-to-cell assignments. The variable names are again the gRNA IDs:

```
gasperini_data['guide_assignment'].var_names[:5].tolist()
```

```
## ['grna_CCGGGCG', 'grna_TGGCGGC', 'grna_AAGGCCG', 'grna_GACGCCG', 'grna_CACACCC']
```

There are no required **vars** for this modality, because the relevant metadata are already in the **guide** modality.

## 2.4 Other fields

There are two other fields of the **gasperini\_data** object, **obs** and **uns**. The **obs** field contains cell metadata not specific to any modality. The most important such metadata is **batch**.

```
gasperini_data.obs
```

```
##          batch
## cell_1      b1
## cell_2      b1
## cell_3      b1
## cell_4      b1
## cell_5      b1
## ...        ...
## cell_44304   b2
## cell_44305   b2
## cell_44306   b2
## cell_44307   b2
## cell_44308   b2
##
## [44308 rows x 1 columns]
```

I propose we have a required variable called **batch** that specifies the batch for each cell. Even if the data only has one batch, we can have a variable with just one value. The other field is **uns**, which contains unstructured metadata. I have included two fields in **uns**: **moi** and **inference\_results**. The **moi** field specifies the MOI of the experiment (high or low):

```
gasperini_data.uns['moi'][0]
```

```
## 'high'
```

I propose for this to be a mandatory field. The `inference_results` field is where the results of the inference will be stored:

```
pd.set_option('display.max_columns', None)
pd.DataFrame(gasperini_data.uns['inference_results'])
```

```
##           gene_id      grna_target  log2_FC  p_value      pair_type
## 0  ENSG00000069482  ENSG00000069482    -9.0    -9.0  positive_control
## 1  ENSG00000100316  ENSG00000100316    -9.0    -9.0  positive_control
## 2  ENSG00000104131  ENSG00000104131    -9.0    -9.0  positive_control
## 3  ENSG00000122026  ENSG00000122026    -9.0    -9.0  positive_control
## 4  ENSG00000135821  ENSG00000135821    -9.0    -9.0  positive_control
## ..           ...           ...           ...           ...           ...
## 615 ENSG00000131094  candidate_enh_8    -9.0    -9.0      discovery
## 616 ENSG00000136448  candidate_enh_8    -9.0    -9.0      discovery
## 617 ENSG00000172992  candidate_enh_8    -9.0    -9.0      discovery
## 618 ENSG00000181513  candidate_enh_8    -9.0    -9.0      discovery
## 619 ENSG00000161714  candidate_enh_8    -9.0    -9.0      discovery
##
## [620 rows x 5 columns]
```

The columns `gene_id` and `grna_target` specify what pairs of gene and targeted element to test for association. The columns `log2_FC` and `p_value` will be filled in by the inference method, and should be initialized with placeholders such as -9. Finally, the optional column `pair_type` specifies the type of pair being tested, e.g. positive control or discovery pair.

### 3 Items for discussion

1. We need to work together to settle on the precise specification of the `MuData` format, not just for the inference task but also for upstream tasks like gRNA assignment. We might want to write Python and/or R functions that check whether a given `MuData` object conforms to whatever specification we end up deciding on.
2. Do we agree that we should test for association between *targeted elements* and gene expression rather than between *individual gRNAs* and gene expression? Should we also have an option to test the latter? I think the latter is adding unnecessary complication at this stage.
3. The current `MuData` structure does not include any cell-wise covariates or QC metrics. We have discussed that there are certain cell-wise covariates that most methods would want to use, like library size. However, adding covariates to the `MuData` will require us to standardize more field names, and different methods might want to use different covariates. Therefore, I thought we could just leave it up to the individual methods to compute whichever covariates they require.